




# Processes and Threads



# Process in Memory

A process is a program in execution

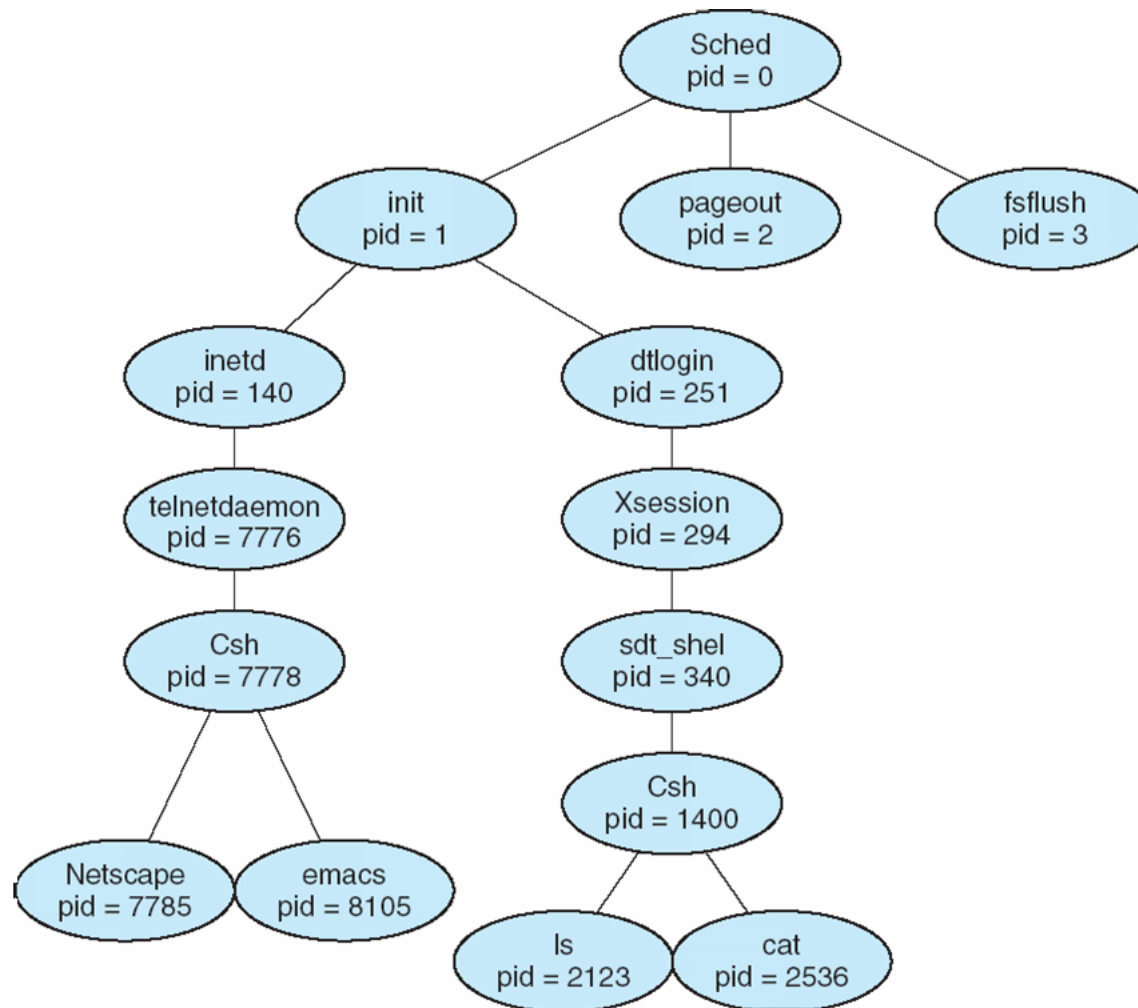


# Multiple processes

---

- Operating systems (OSes) run multiple processes on behalf of each user (example on next slide)
- Processes have a parent-child relationship
- Users can also write programs that use multiple processes – it is one kind of parallel programming
- Key function calls: `fork()`, family of `exec()` functions
- Fork and exec are system calls that each language provides a way to access
- It is easy to do in C/C++ since OSes are written in C

# A tree of processes on a typical Solaris



# C Program forking separate process

---

```
int main()
{
    pid_t  pid;
    /* fork another process */
    pid = fork();
    if (pid < 0) { /* error occurred */
        fprintf(stderr, "Fork Failed");
        exit(-1);
    }
    else if (pid == 0) { /* child process */
        execlp("/bin/ls", "ls", NULL);
    }
    else { /* parent process */
        /* parent will wait for the child to
        complete */
        wait (NULL);
        printf ("Child Complete");
        exit(0);
    }
}
```

# Why bother with parallel programming?

---

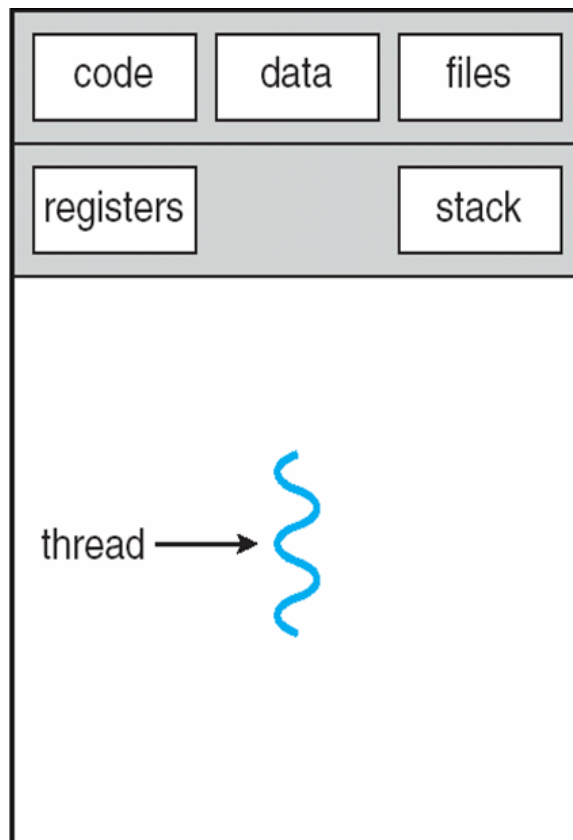
- Modern machines are multicore and parallel programming can utilize the idle cores
- If processing must stop due to network or disk I/O reasons, parallel programming allows other tasks to be completed in the meanwhile

# Threads

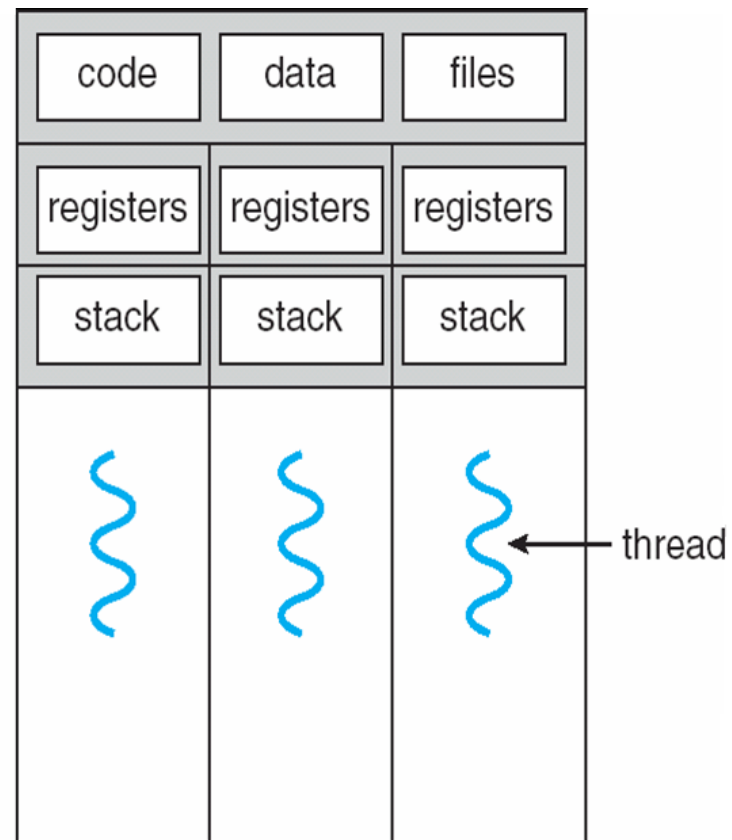
---

- A fundamental unit of CPU utilization
- A thread is a baby process
- Yet another way of doing parallel programming
- Choice between using process vs thread is task specific
- Threads are lower overhead
- Threads share things (see next slide)

# Single and Multithreaded Processes



single-threaded process



multithreaded process



# Thread synchronization: Mutexes

---

- Mutex = mutual exclusion
- Mutexes are used for serializing shared resources such as memory
- Mutexes can be applied only to threads in a single process and do not work between processes
- Other primitives, such as, semaphores, exist for that purpose but outside the scope of this class

## Threaded function without mutex

---

```
int counter = 0;  
void functionC()  
{  
    counter++;  
}
```

■ Possible execution sequence for two threads:

Thread 1: counter = 0 → counter = 1;

Thread 2: counter = 0 → counter = 1;

■ Undesirable, we need counter = 2;

## Threaded function with mutex

---

```
pthread_mutex_t mutex1 = PTHREAD_MUTEX_INITIALIZER;  
int counter = 0;  
void functionC() {  
    pthread_mutex_lock(&mutex1);  
    counter++;  
    pthread_mutex_unlock(&mutex1);}
```

■ Possible execution sequence for two threads:

Thread 1: counter = 0 → counter = 1;

Thread 2: counter = 0 → lockout → counter = 1;

■ Threads 1 and 2 may be swapped in incrementing counter but counter = 2 when mutex is used

## Pthread\_mutex\_trylock()

---

- What if multiple mutexes are in use and different threads end up locking different mutexes, potentially blocking each other?
- Such conditions are called deadlock conditions
- pthread\_mutex\_lock() blocks a thread until mutex is unlocked
- pthread\_mutex\_trylock() prevents deadlock by returning an error code to signal failure to acquire a lock

## Thread pitfalls: race conditions

---

- Threads are scheduled by the operating system and are executed at random
- It cannot be assumed that threads are executed in the order they are created
- Threads may also execute at different speeds
- Mutexes and joins must be utilized to achieve a predictable execution order and outcome

## Thread pitfalls: thread safe code

---

- The threaded routines must call functions which are "thread safe"
- This means that there are no static or global variables which other threads may clobber or read assuming single threaded operation
- If static or global variables are used then mutexes must be applied or the functions must be re-written to avoid the use of these variables

## Thread pitfalls: mutex deadlock

---

- This condition occurs when a mutex is applied but then not "unlocked"
- This causes program execution to halt indefinitely
- It can also be caused by poor application of mutexes or joins
- Be careful when applying two or more mutexes to a section of code